



oci-dev-binder-hook

dynamic device access management based on container annotations

Roberto Majadas <rmajadas@redhat.com>
Albert Esteve <aesteve@redhat.com>



The Challenge

- Containers currently offer static device addition using `--device`
- But if you're working with diverse hardware boards (e.g., automotive or IoT devices), the exact set of peripheral devices is often unknown beforehand.
 - For instance, the [QM container](#) could require access to devices that could be represented in different way depending on the board, i.e: `/dev/dri/cardX /dev/input/*`



What can help us to solve the challenge?

- The **Open Container Initiative** (OCI) standard defines a runtime specification for containers.
 - This standard includes **hooks**, which are scripts or programs that can be triggered at different points in the container lifecycle.
 - We can use a **pre-start** hook to run a program **after** the container namespace is created but **before** the user's application starts.
- **Systemd/udev**: The standard Linux subsystem for dynamic device detection and management. It allows us to query for devices based on their properties and tags.
- [Multi-Seat on Linux](#) from Freedesktop:
 - Defines seats as a set of devices integrated with systemd/udev



oci-dev-binder-hook

- An oci-hook that inspects the udev database looking for the devices that you want to include.
- The hook can be configured using container annotations
 - `podman run --annotation io.dev-binder.udev.seat=seat0 -it my-application:latest`
- We can let the dynamic device management to udev and just retrieve the tagged devices.



Demo

```
oci-dev-binder-hook on └─ main [??] is  v0.1.0 via  v1.89.0  
• > podman run --annotation=io.dev-binder.udev.seat=seat0 --rm -it fedora find /dev/dri/card*  
/dev/dri/card1
```



Next steps (?)

- Move the project to github.com/containers
- Improve the annotations' semantics
 - Filter by subsystem tag (?)



Thanks!
Questions?